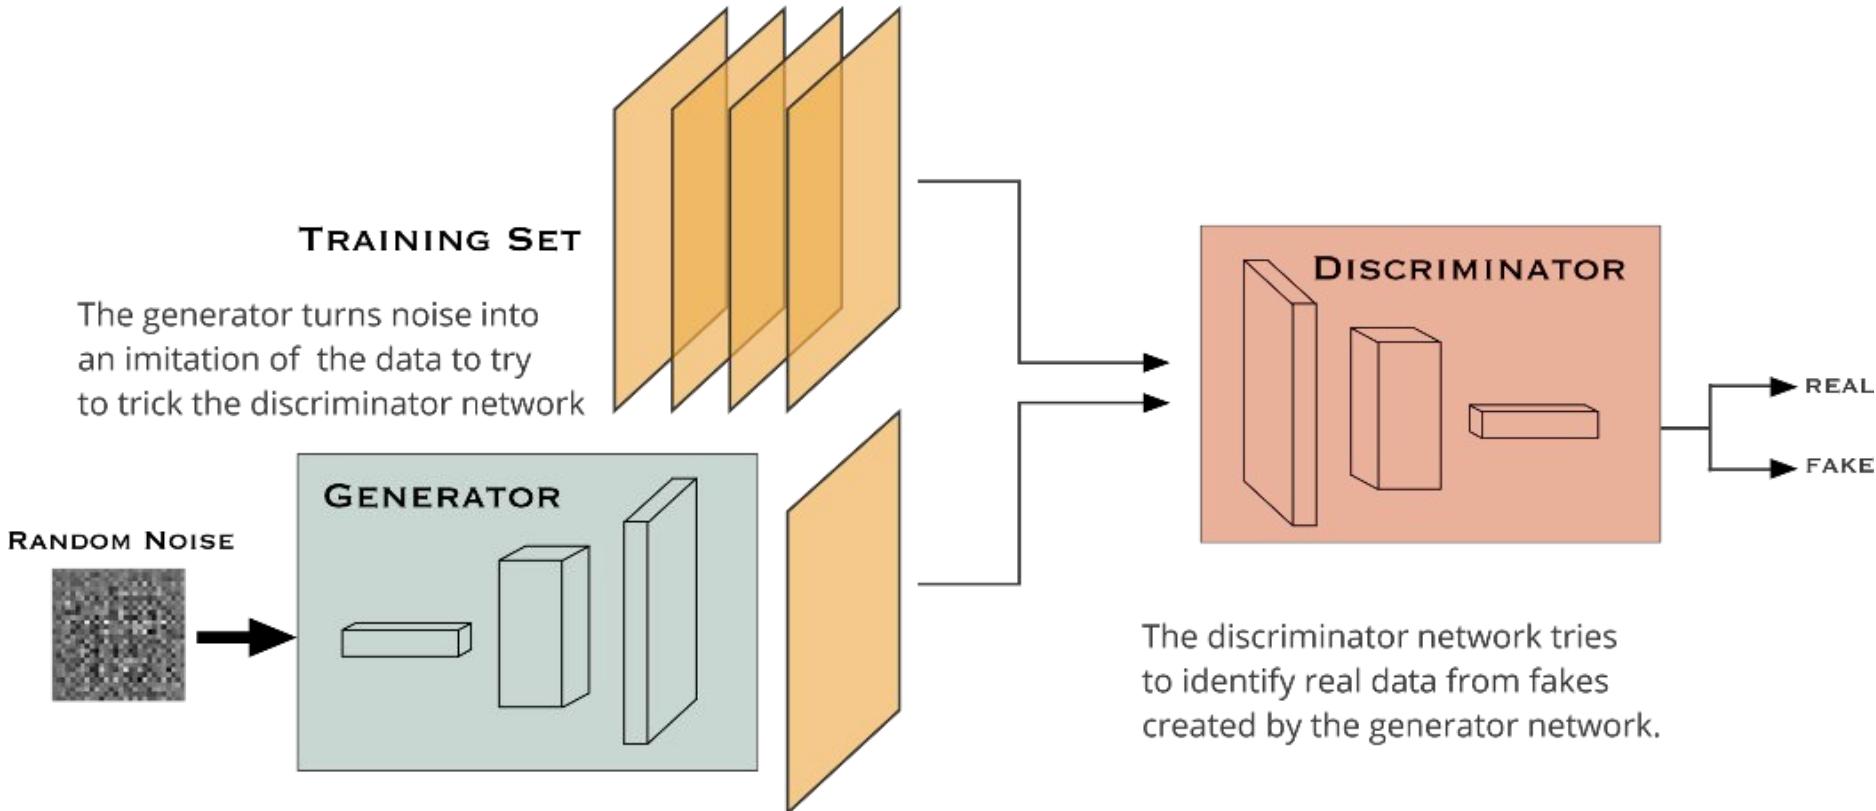


Exploring Image Generation

Harth Majeed
3546 Deep Learning

Quick Overview

We “generate” images to trick the “discriminator”.



How does a GAN react to different types of Images?

- How does the network react to different types of images
- Does it react well to textures?
- Does it react well to shapes?
- Does it react well to detail?

We wish to understand the strength and weaknesses

Baseline Model - Generator

- Sequential()
- model.add(layers.Dense(8*8*512, use_bias=False))
- model.add(layers.Reshape((8, 8, 512)))
- assert model.output_shape == (None, 8, 8, 512)
- model.add(layers.Conv2DTranspose(256, (5, 5), strides=(1, 1), padding='same', use_bias=False))
- assert model.output_shape == (None, 8, 8, 256)
- model.add(layers.Conv2DTranspose(128, (5, 5), strides=(2, 2), padding='same', use_bias=False))
- assert model.output_shape == (None, 16, 16, 128)
- model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', use_bias=False))
- assert model.output_shape == (None, 32, 32, 64)
- model.add(layers.Conv2DTranspose(32, (5, 5), strides=(2, 2), padding='same', use_bias=False))
- assert model.output_shape == (None, 64, 64, 32)
- model.add(layers.Conv2DTranspose(16, (5, 5), strides=(2, 2), padding='same', use_bias=False))
- assert model.output_shape == (None, 128, 128, 16)
- model.add(layers.Conv2DTranspose(3, (5, 5), strides=(2, 2), padding='same', use_bias=False, activation='tanh'))
- assert model.output_shape == (None, 256, 256, 3)

Baseline Model - Discriminator

- Sequential()
- model.add(layers.Conv2D(16, (5, 5), strides=(2, 2), padding='same', input_shape=[256, 256, 3]))
- model.add(layers.LeakyReLU())
- model.add(layers.Dropout(0.3))
- model.add(layers.Conv2D(32, (5, 5), strides=(2, 2), padding='same'))
- model.add(layers.LeakyReLU())
- model.add(layers.Dropout(0.3))
- model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same'))
- model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
- model.add(layers.Conv2D(256, (5, 5), strides=(2, 2), padding='same'))
- model.add(layers.Flatten())
- model.add(layers.Dense(1))

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32768)	3276800
batch_normalization (BatchN ormalization)	(None, 32768)	131072
leaky_re_lu (LeakyReLU)	(None, 32768)	0
reshape (Reshape)	(None, 8, 8, 512)	0
conv2d_transpose (Conv2DTra nspose)	(None, 8, 8, 256)	3276800
batch_normalization_1 (Bac hNormalization)	(None, 8, 8, 256)	1024
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 256)	0
conv2d_transpose_1 (Conv2DT ranspose)	(None, 16, 16, 128)	819200
batch_normalization_2 (Bac hNormalization)	(None, 16, 16, 128)	512
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_2 (Conv2DT ranspose)	(None, 32, 32, 64)	204800
batch_normalization_3 (Bac hNormalization)	(None, 32, 32, 64)	256
leaky_re_lu_3 (LeakyReLU)	(None, 32, 32, 64)	0
conv2d_transpose_3 (Conv2DT ranspose)	(None, 64, 64, 32)	51200
batch_normalization_4 (Bac hNormalization)	(None, 64, 64, 32)	128
leaky_re_lu_4 (LeakyReLU)	(None, 64, 64, 32)	0
conv2d_transpose_4 (Conv2DT ranspose)	(None, 128, 128, 16)	12800
batch_normalization_5 (Bac hNormalization)	(None, 128, 128, 16)	64
leaky_re_lu_5 (LeakyReLU)	(None, 128, 128, 16)	0
conv2d_transpose_5 (Conv2DT ranspose)	(None, 256, 256, 3)	1200
<hr/>		
Total params:	7,175,856	
Trainable params:	7,709,328	
Non-trainable params:	66,528	

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 16)	1216
leaky_re_lu_6 (LeakyReLU)	(None, 128, 128, 16)	0
dropout (Dropout)	(None, 128, 128, 16)	0
conv2d_1 (Conv2D)	(None, 64, 64, 32)	12832
leaky_re_lu_7 (LeakyReLU)	(None, 64, 64, 32)	0
dropout_1 (Dropout)	(None, 64, 64, 32)	0
conv2d_2 (Conv2D)	(None, 32, 32, 64)	51264
leaky_re_lu_8 (LeakyReLU)	(None, 32, 32, 64)	0
dropout_2 (Dropout)	(None, 32, 32, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 128)	204928
leaky_re_lu_9 (LeakyReLU)	(None, 16, 16, 128)	0
dropout_3 (Dropout)	(None, 16, 16, 128)	0
conv2d_4 (Conv2D)	(None, 8, 8, 256)	819456
leaky_re_lu_10 (LeakyReLU)	(None, 8, 8, 256)	0
dropout_4 (Dropout)	(None, 8, 8, 256)	0
flatten (Flatten)	(None, 16384)	0
dense_1 (Dense)	(None, 1)	16385
<hr/>		
Total params: 1,106,081		
Trainable params: 1,106,081		
Non-trainable params: 0		

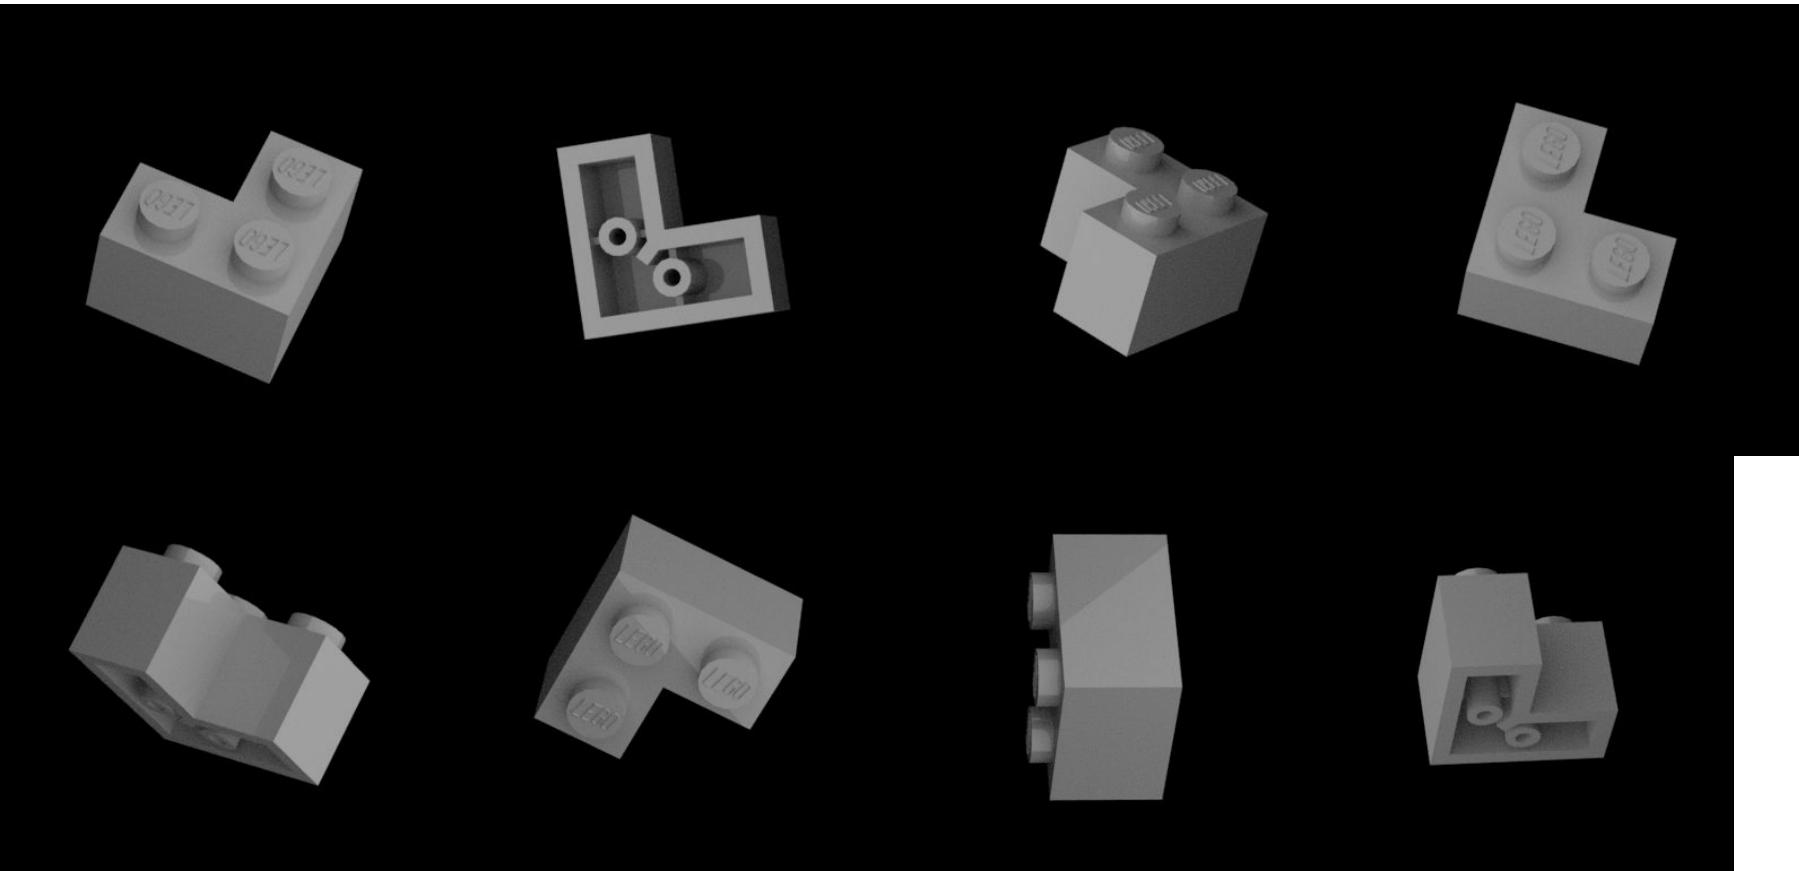
The Plan

- The baseline model will be used across all datasets, will not be changed
- Epoch = 5000

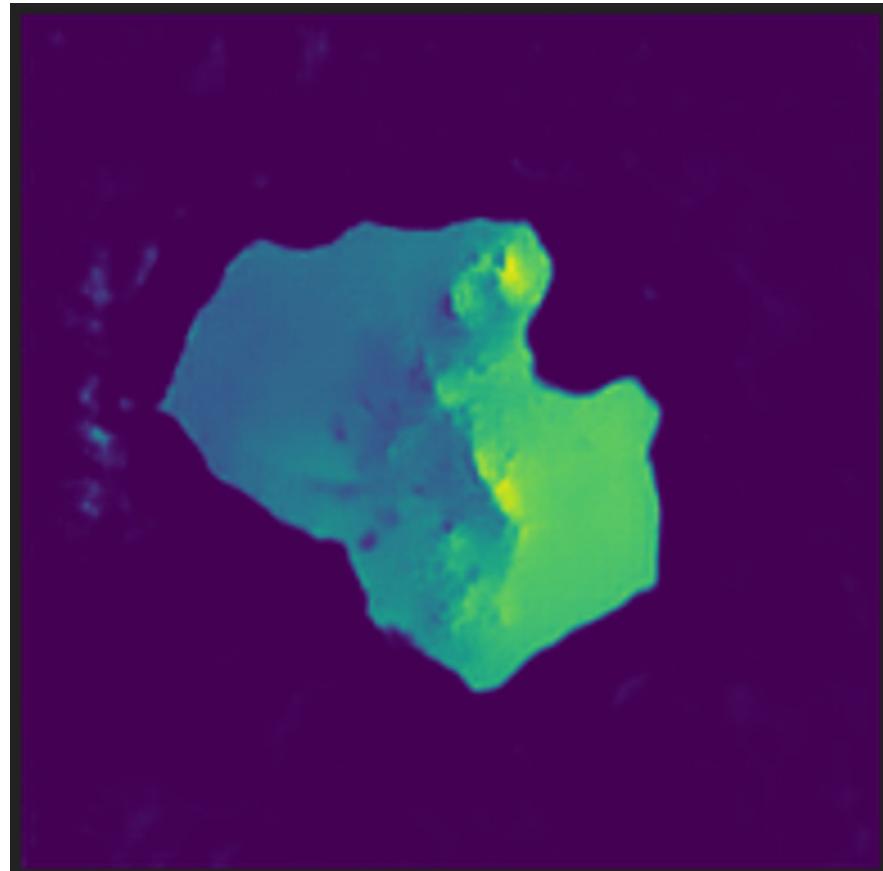
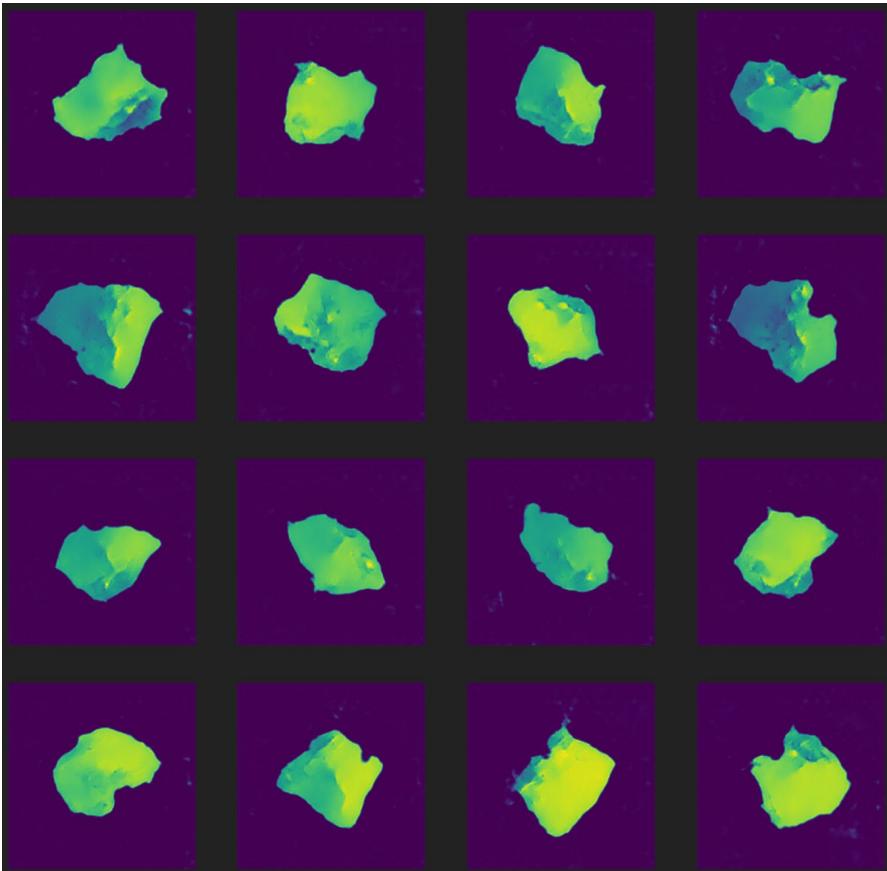
Datasets

- Lego, corner piece (800 images)
- Cutlery (spoons, knifes, forks, etc) (870 images)
- Marble textures (120 images)
- Cats (500, 4096 images)

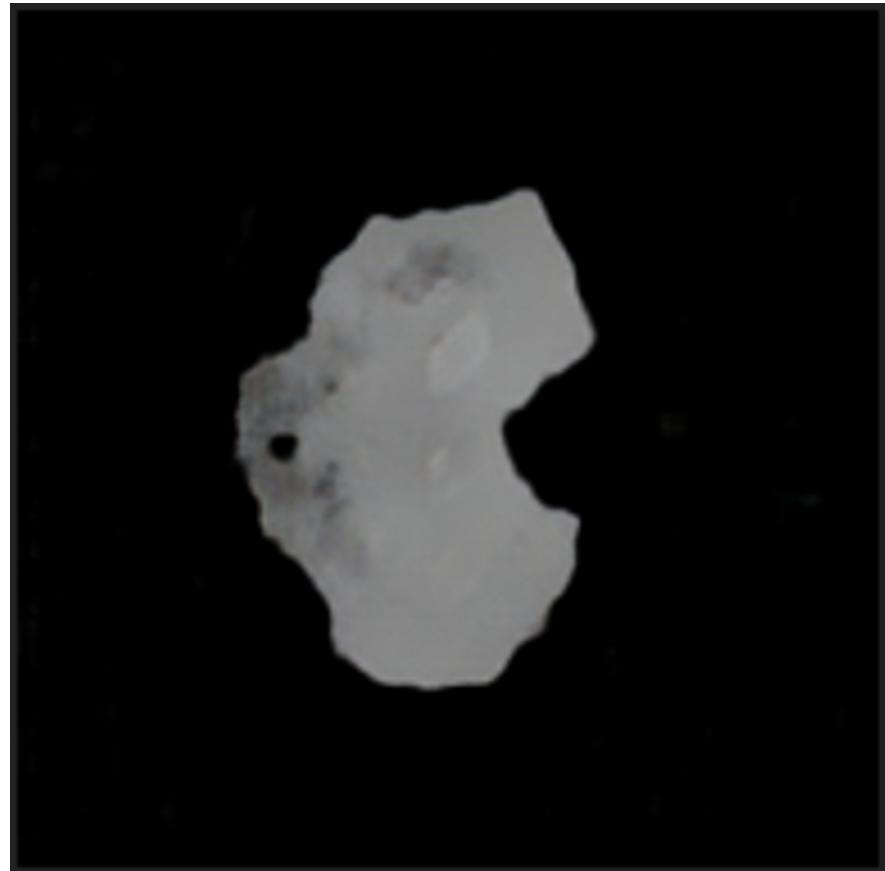
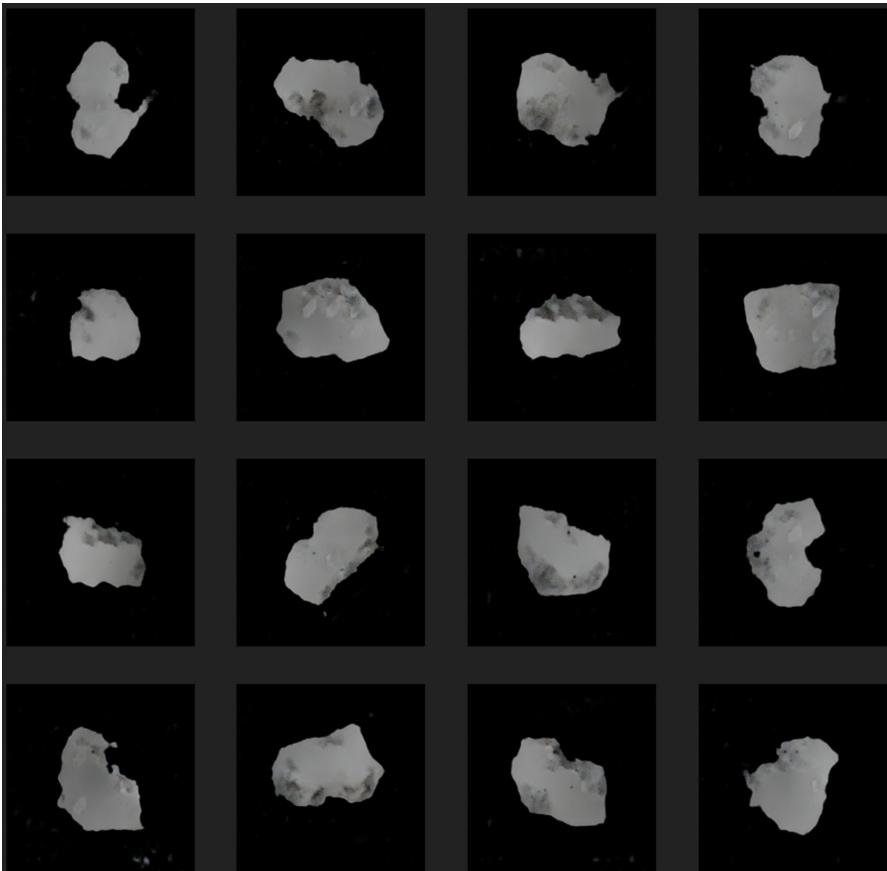
Lego Results - Input Images (800)



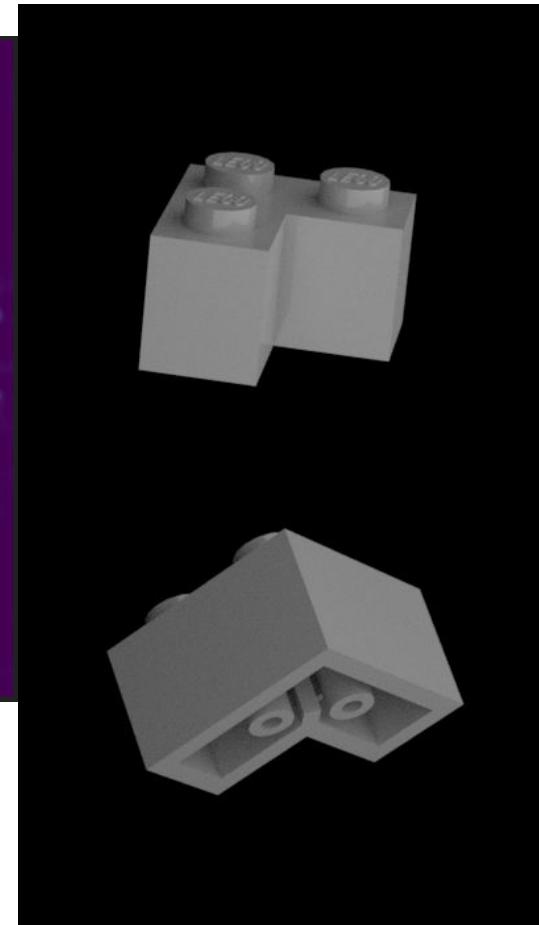
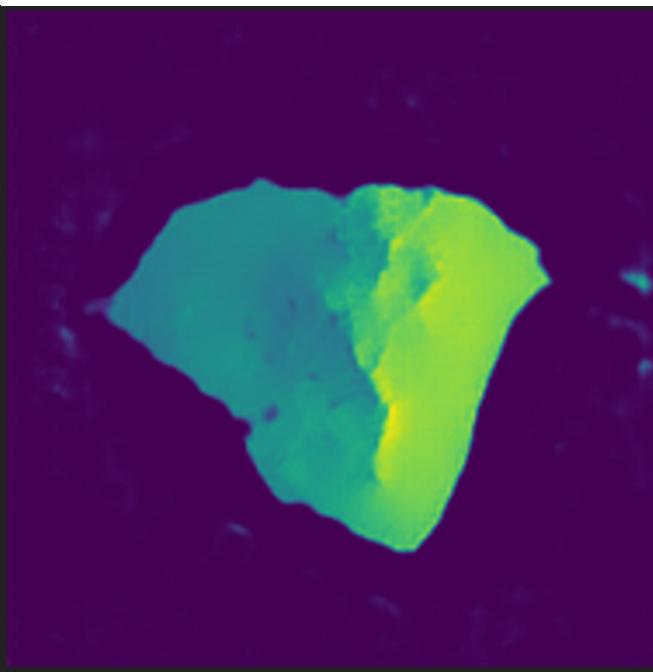
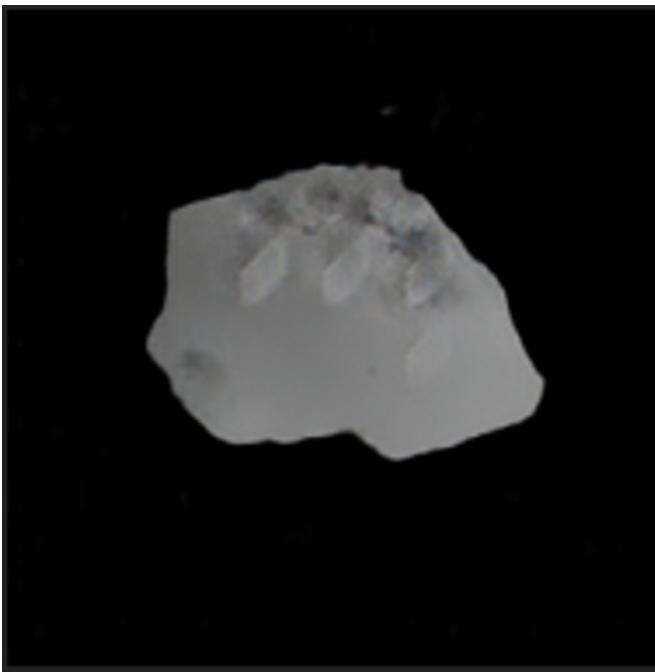
Lego Results - Output Images (800 Images) 5k epoch



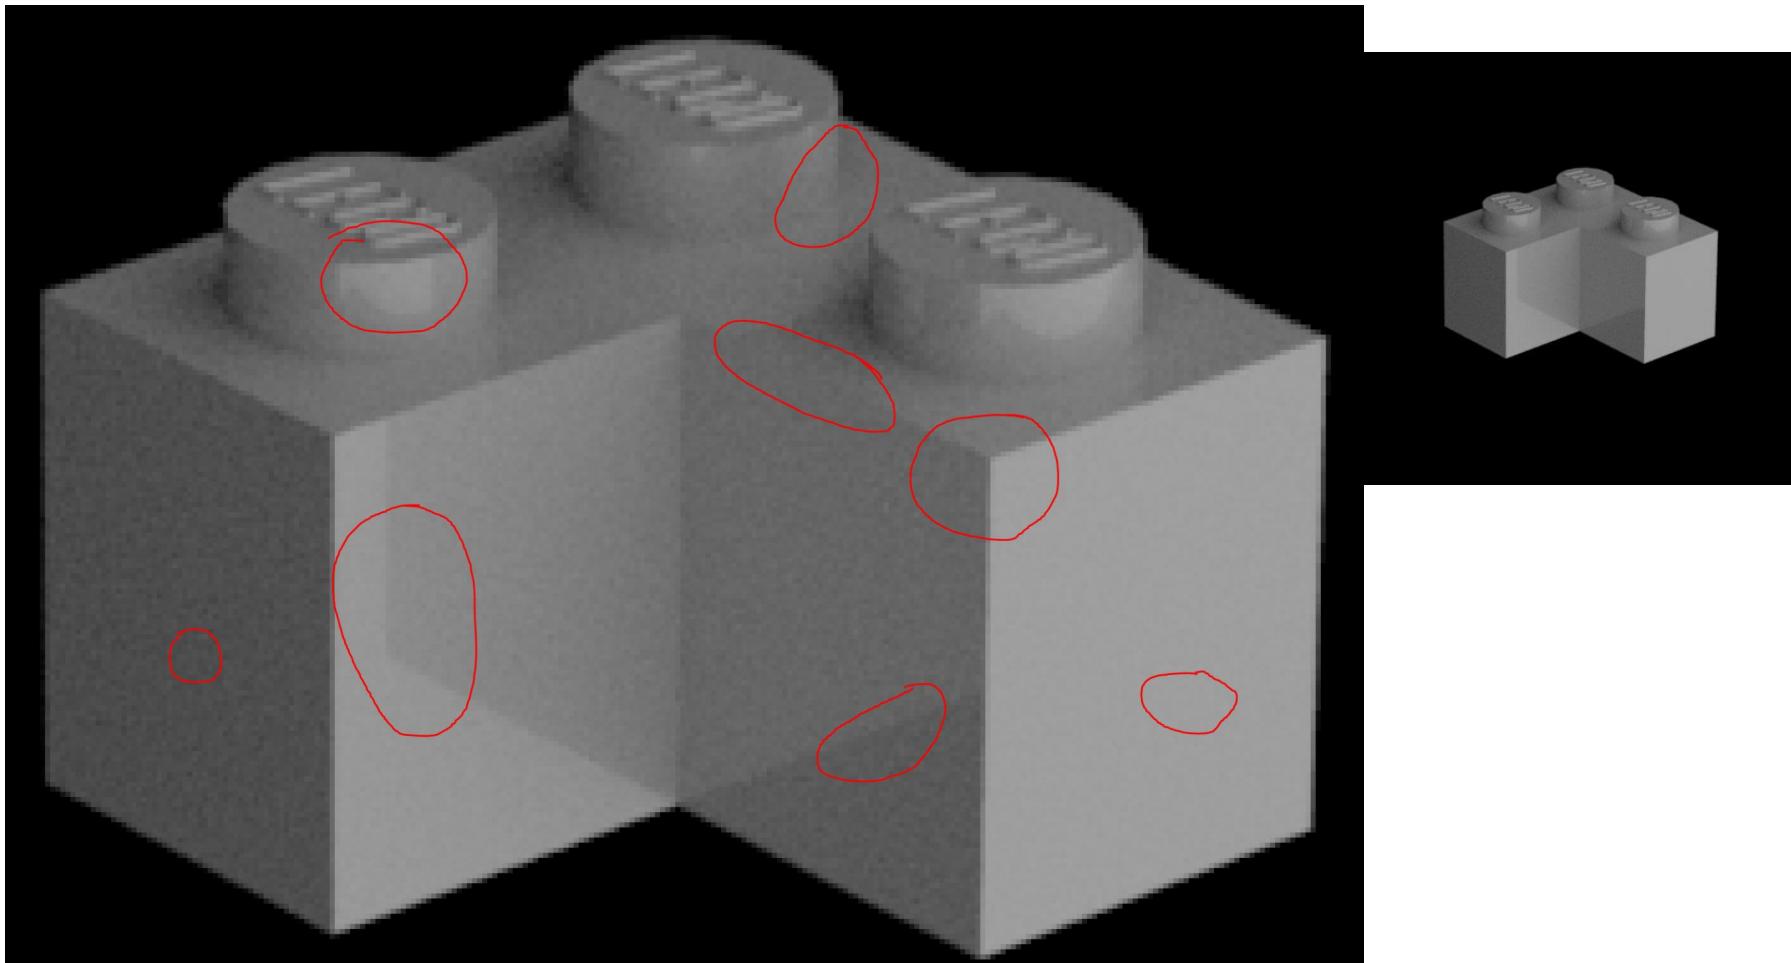
Lego Results - Output Images (800 Images) 5k epoch



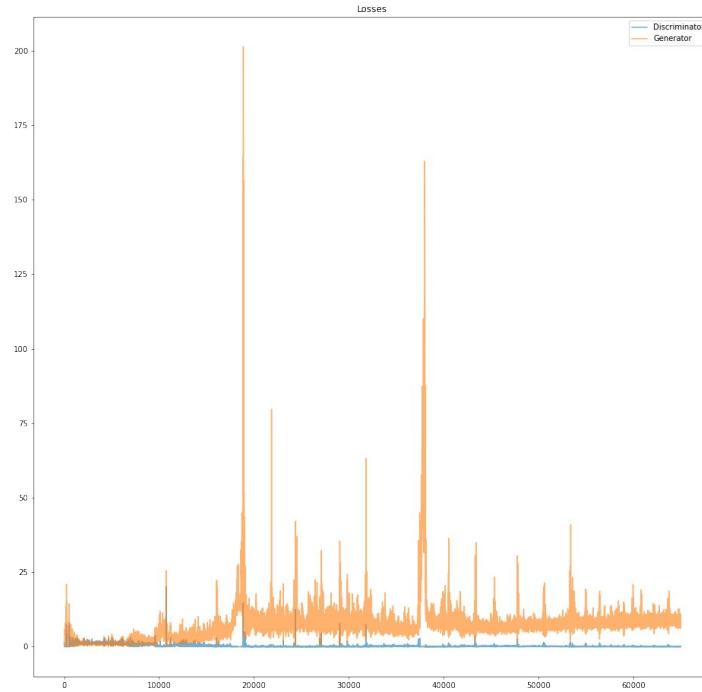
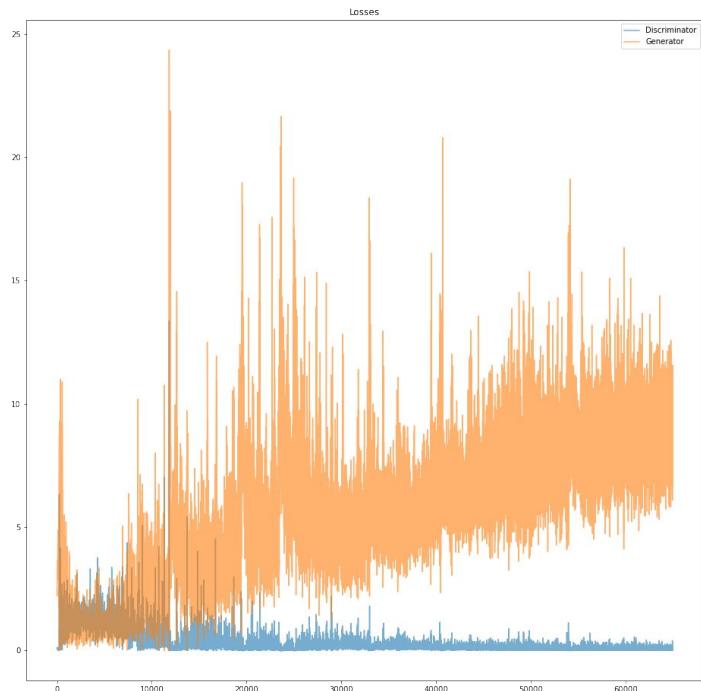
Lego Results - Output Images (800 Images)



Lego Results - Analyzing the image more (800 Images)



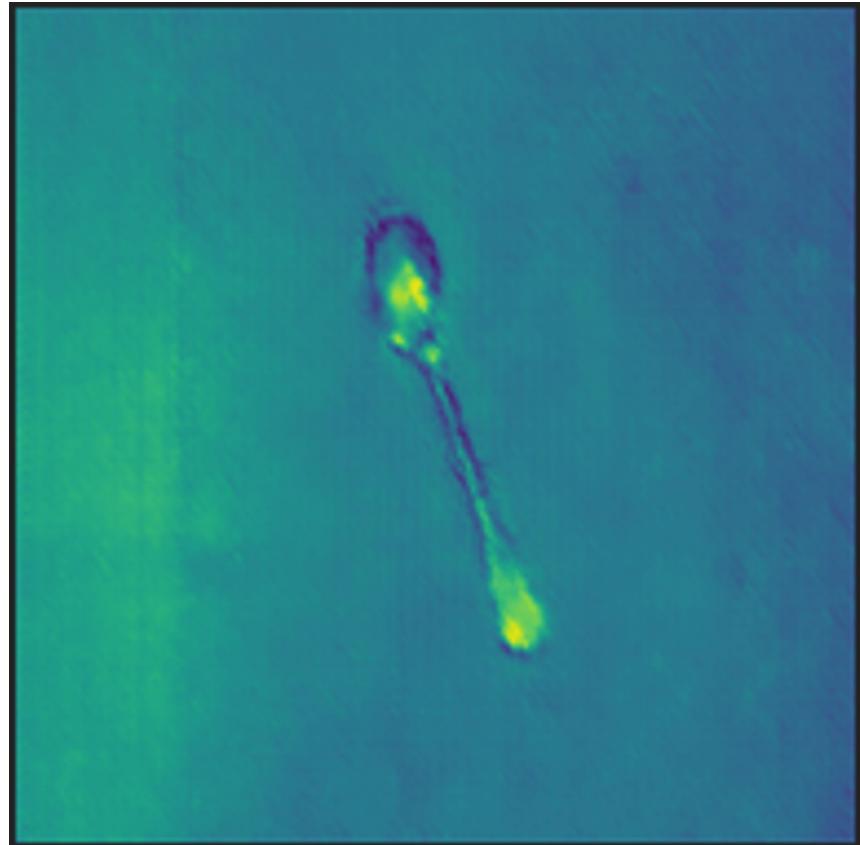
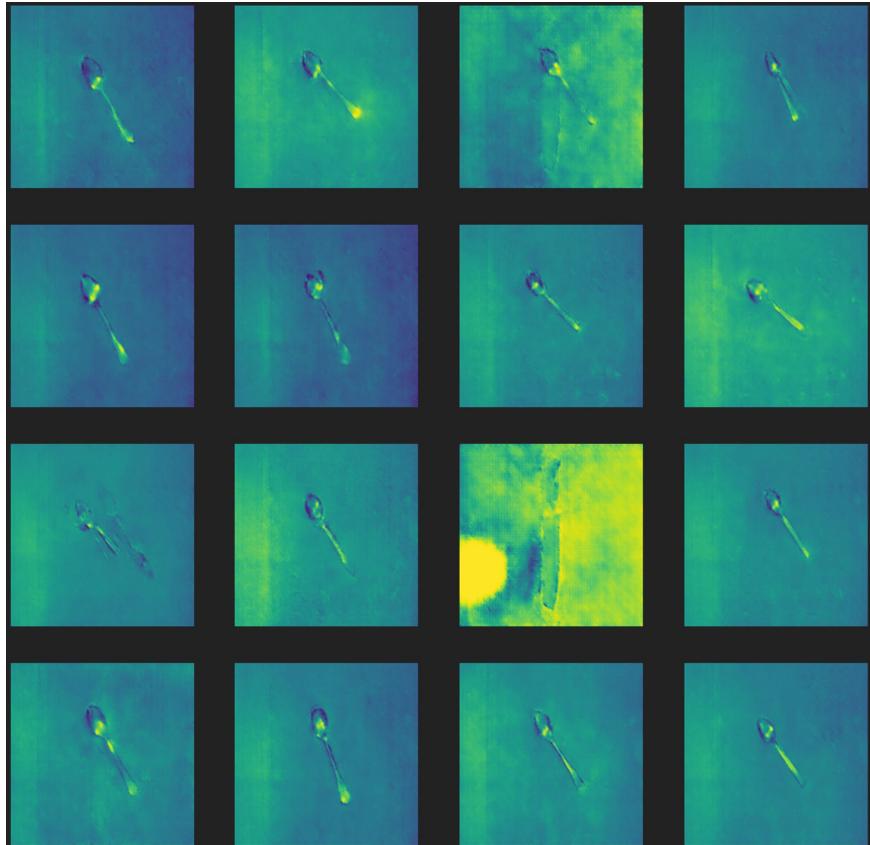
Lego Results - Graphs (800 Images) 5k epoch



Cutlery Results - Input Images (870)



Cutlery Results - Output Images (870) 5k epoch



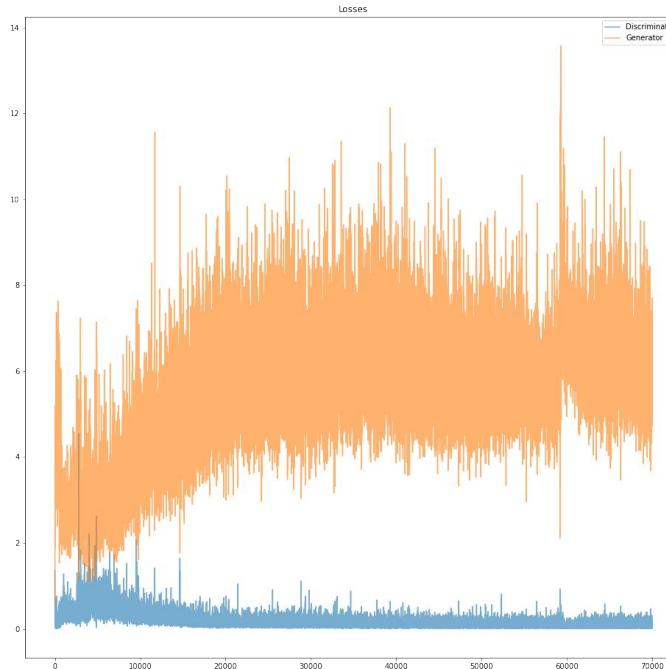
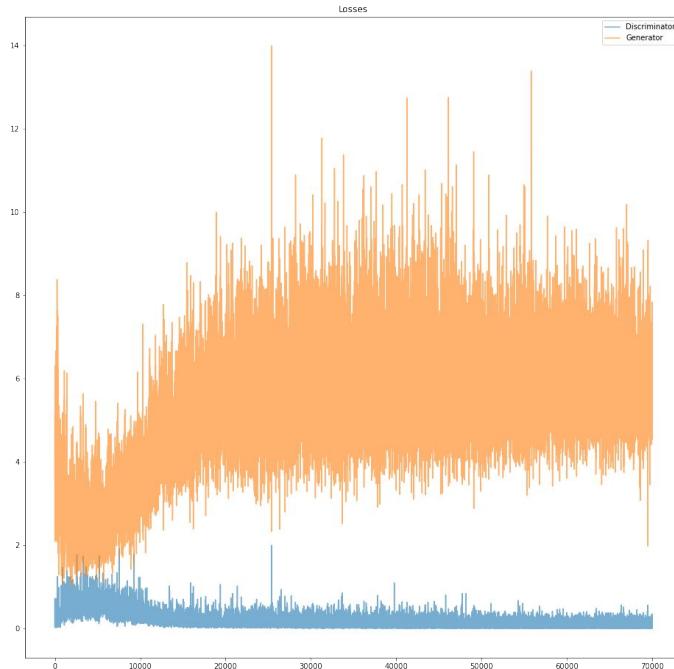
Cutlery Results - Output Images (870) 5k epoch



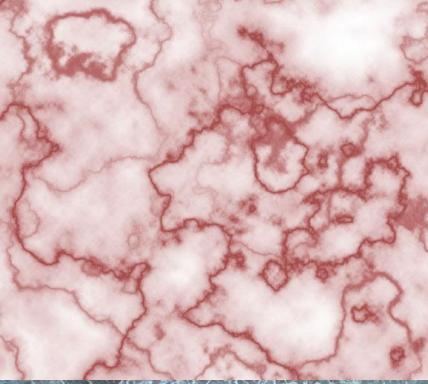
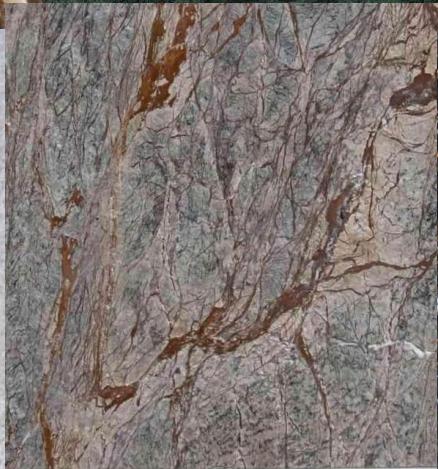
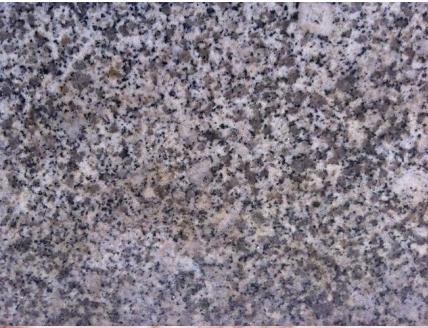
Cutlery Results - Analyzing the images (870)



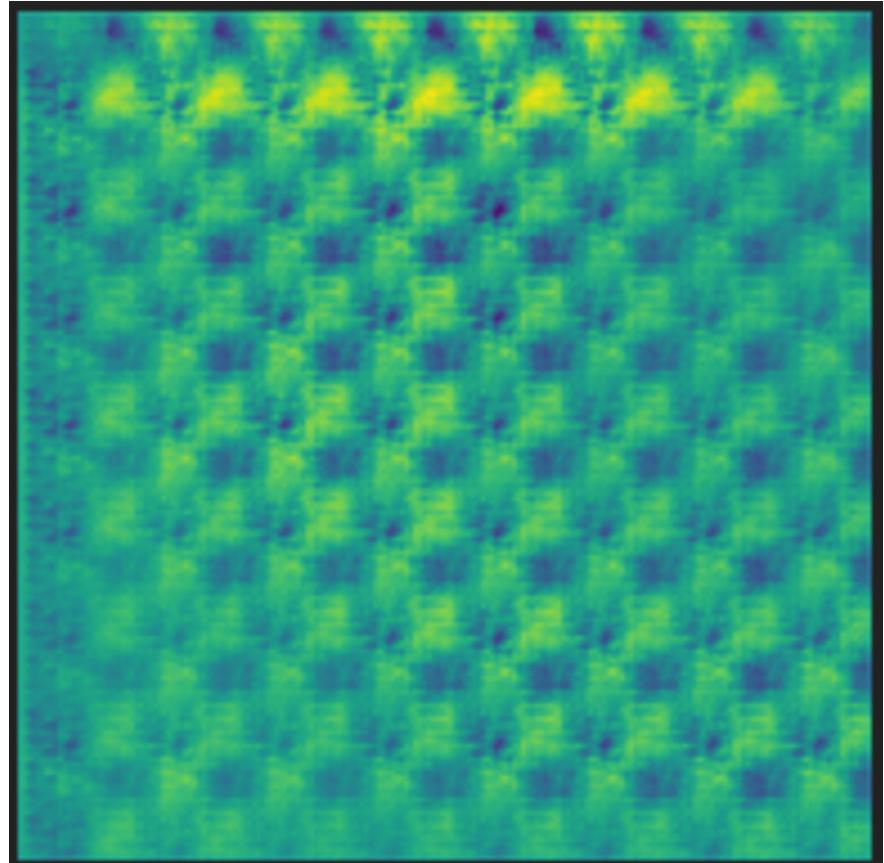
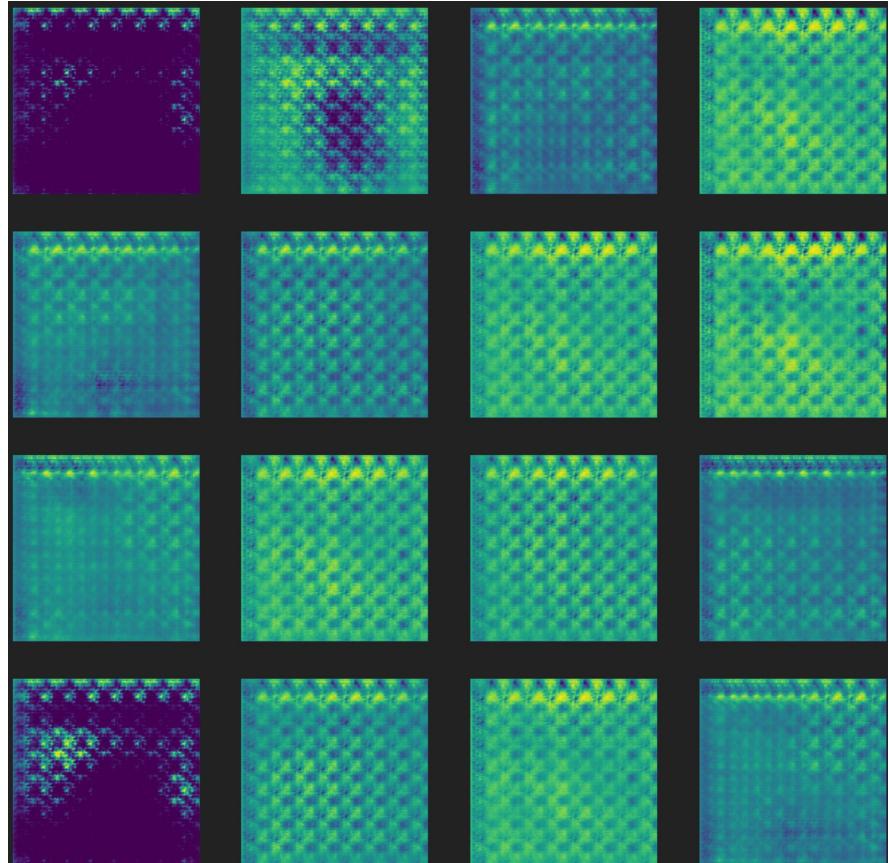
Cutlery Results - Output Images (870) 5k epoch



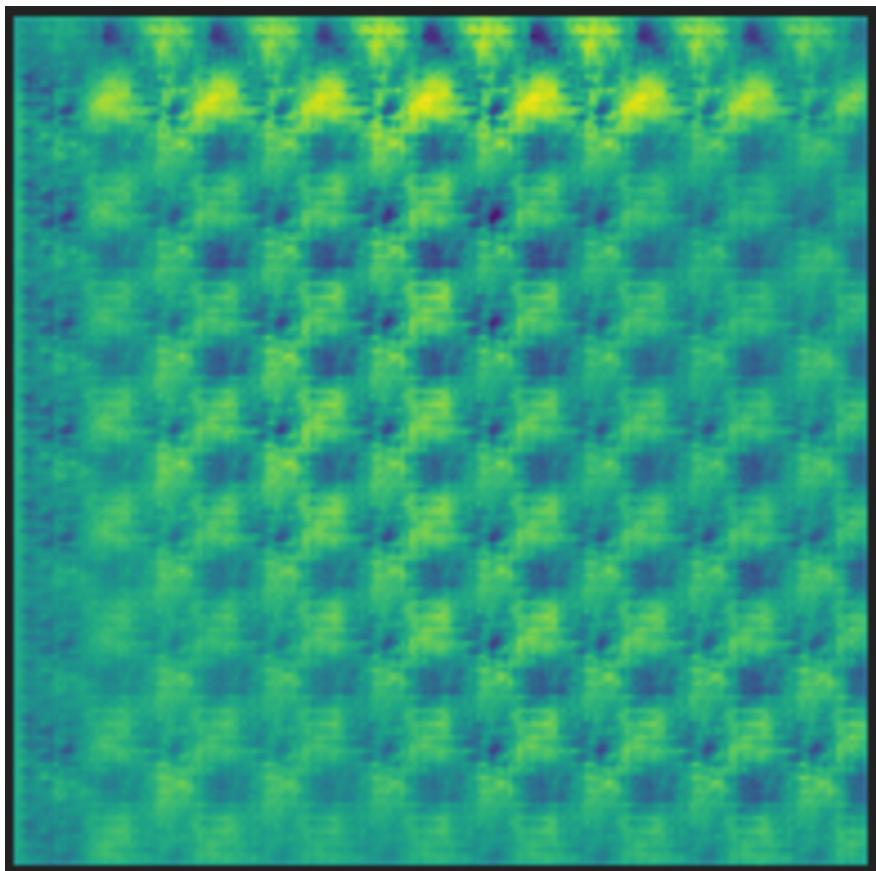
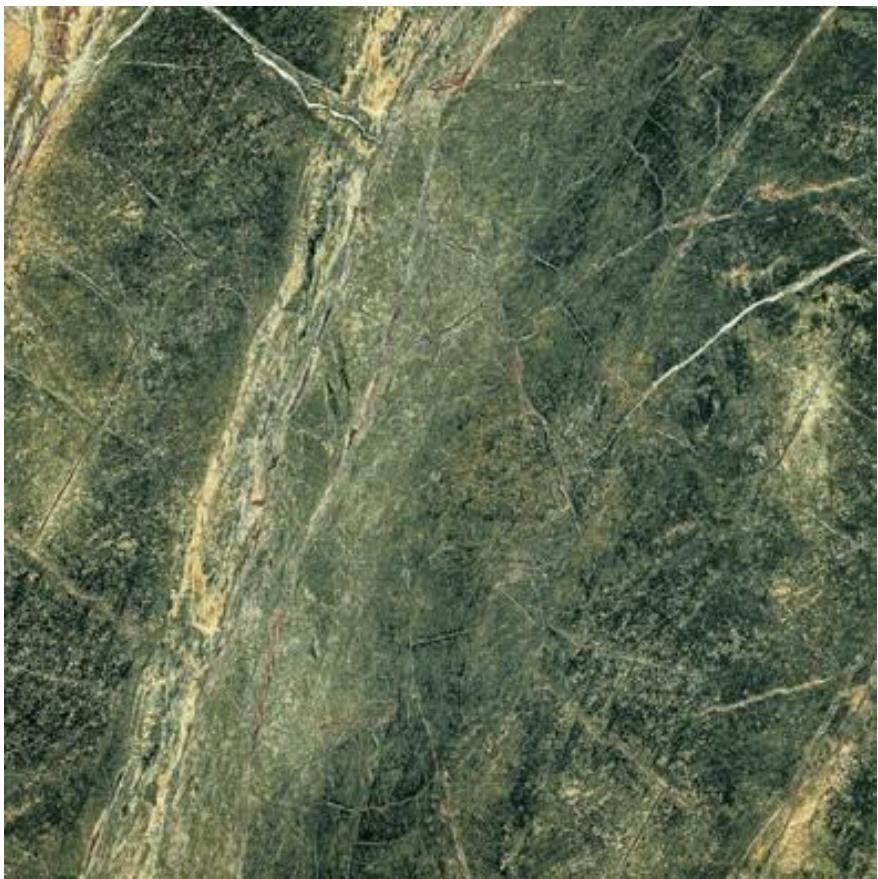
Marble Texture Results - Input Images (120)



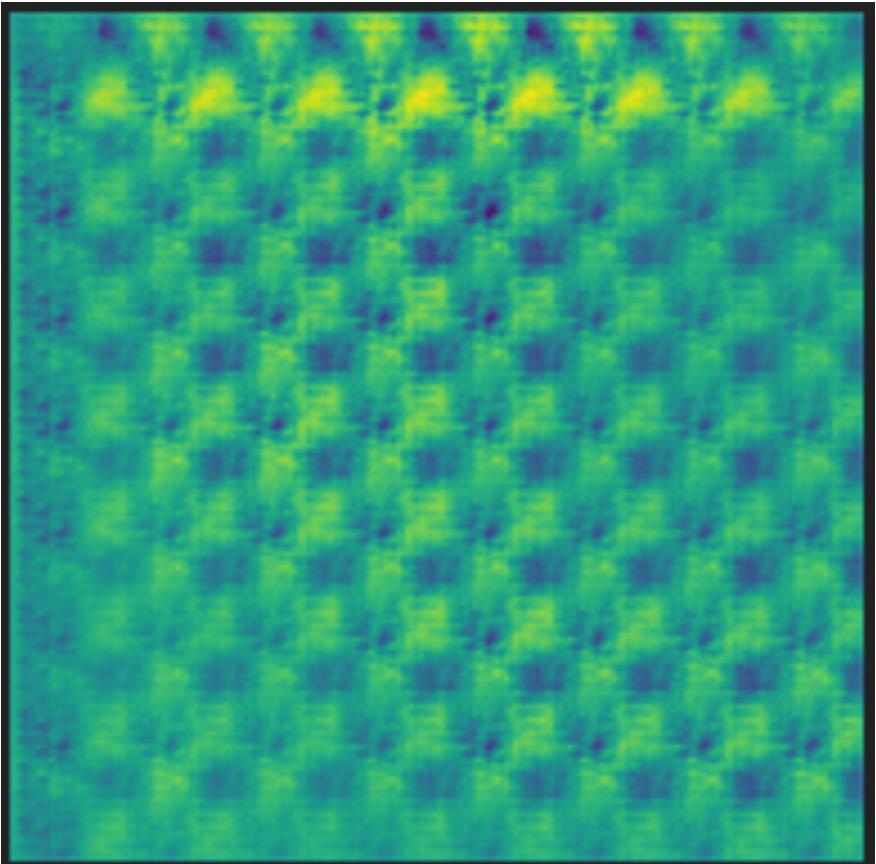
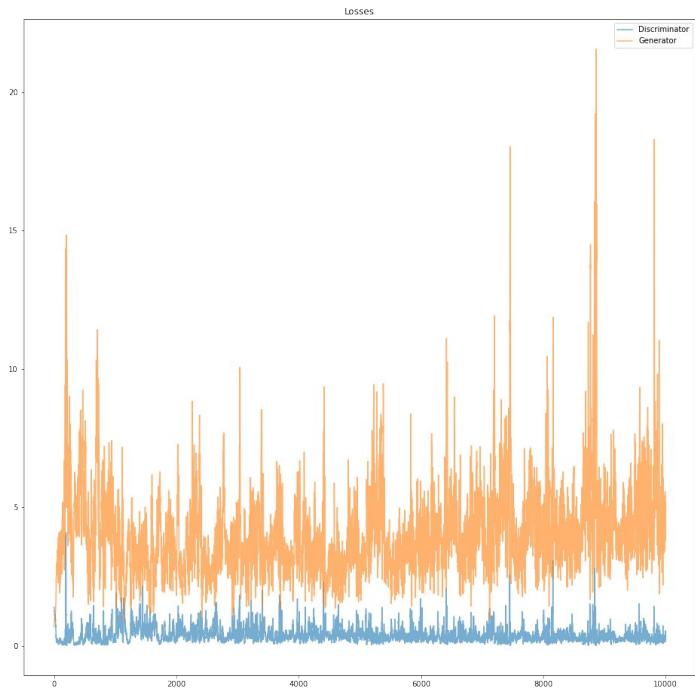
Marble Texture Results - Output Images (120) 5k epoch



Marble Texture Results - Output Images (120) 5k epoch



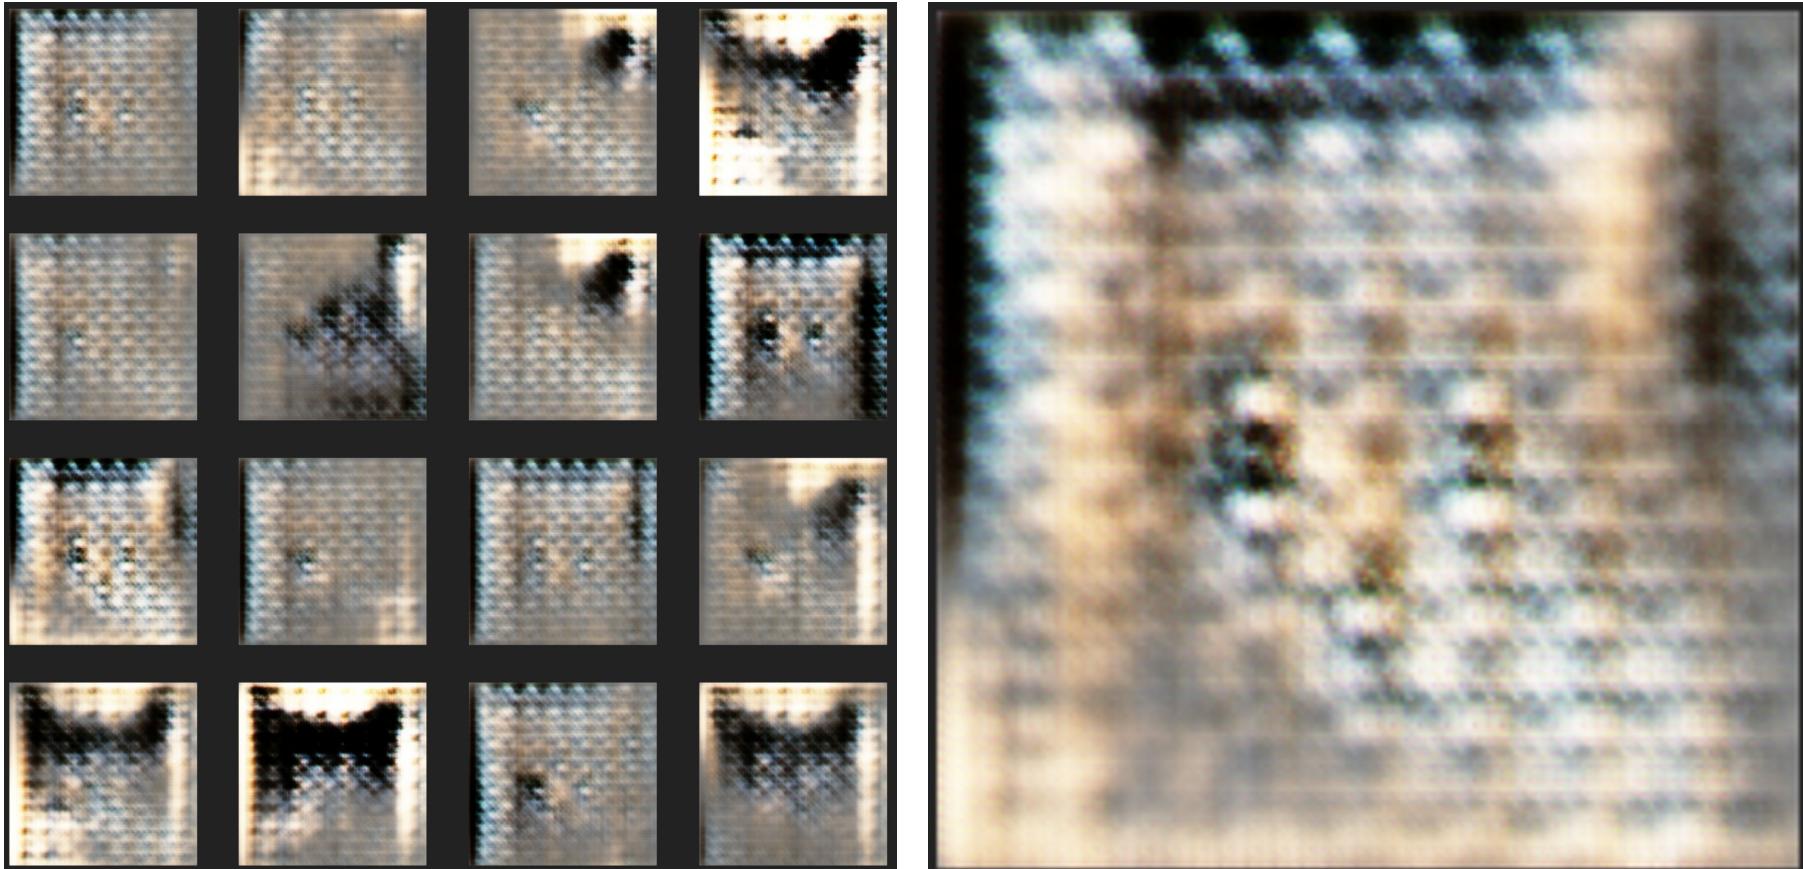
Marble Texture Results - Graphs (120) 5k epoch



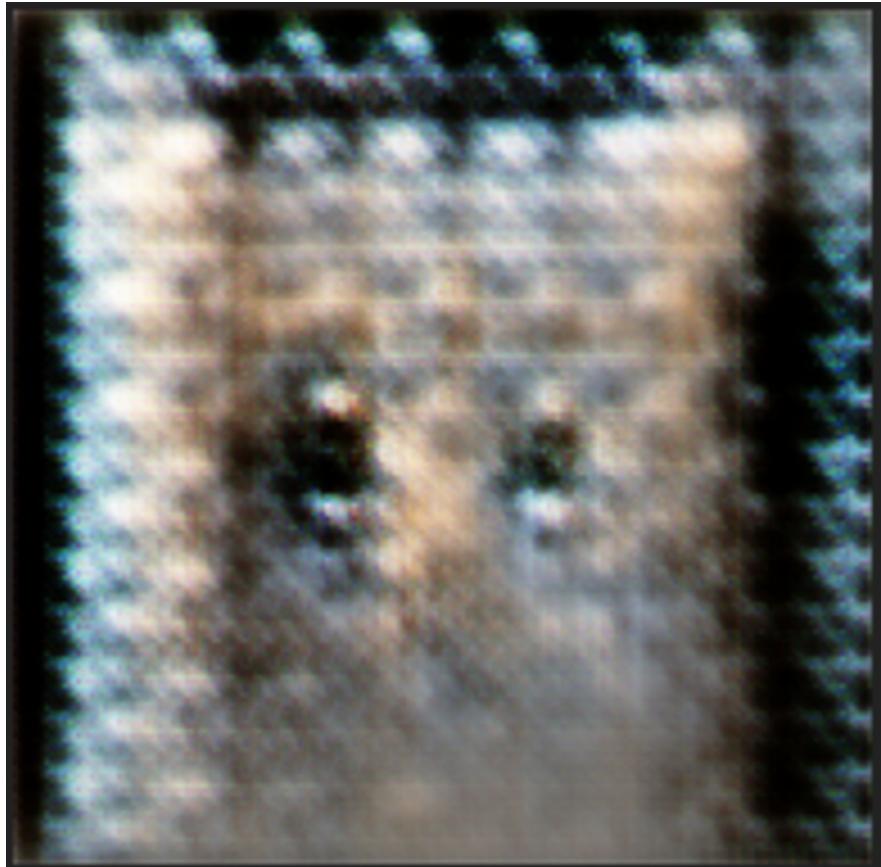
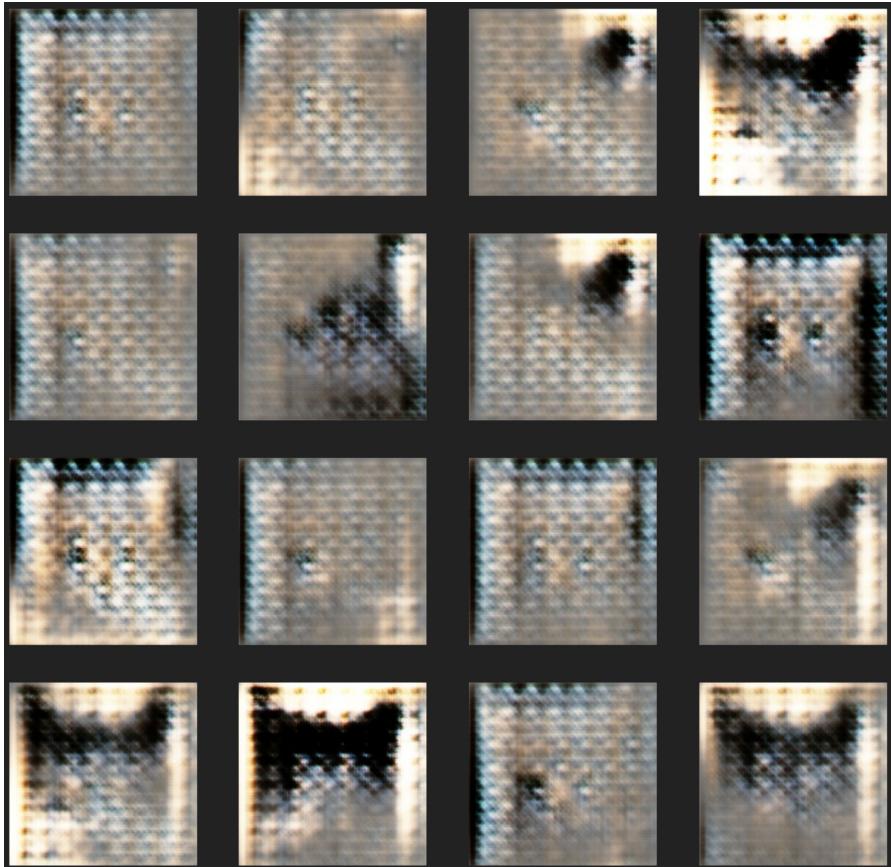
Cats Results - Input Images (500, 4096) 1k epoch



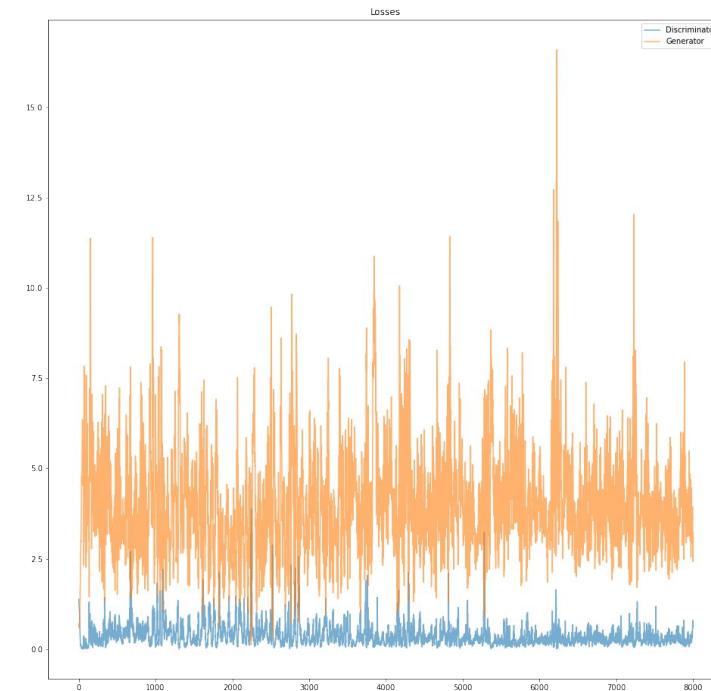
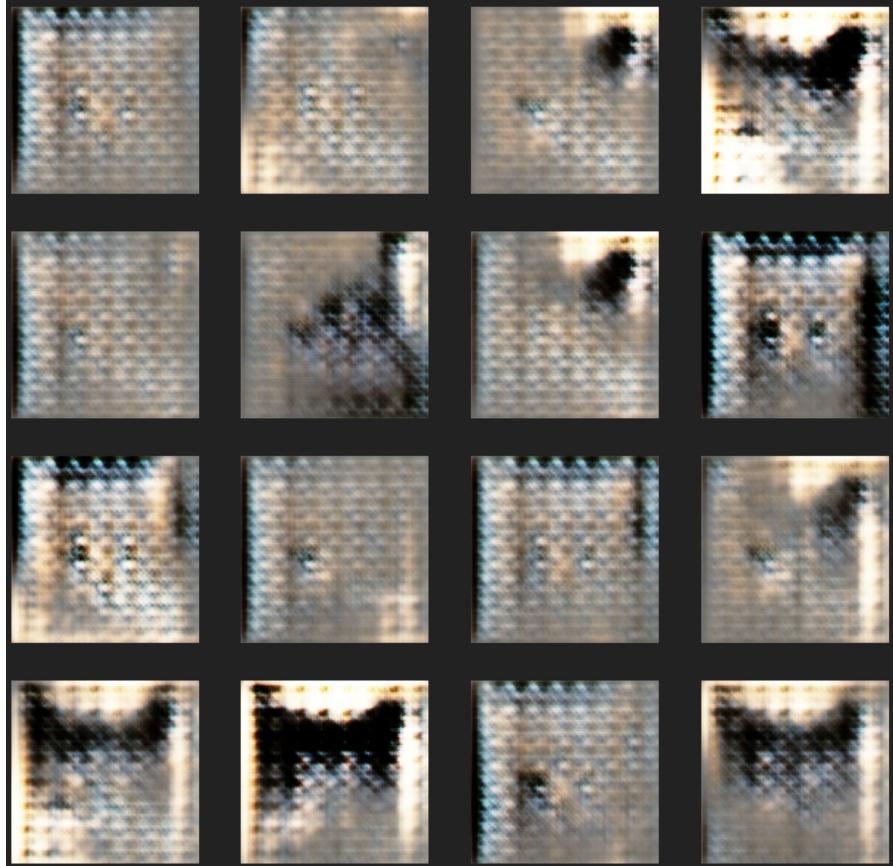
Cats Results - Output Images (500 images) 1k epoch



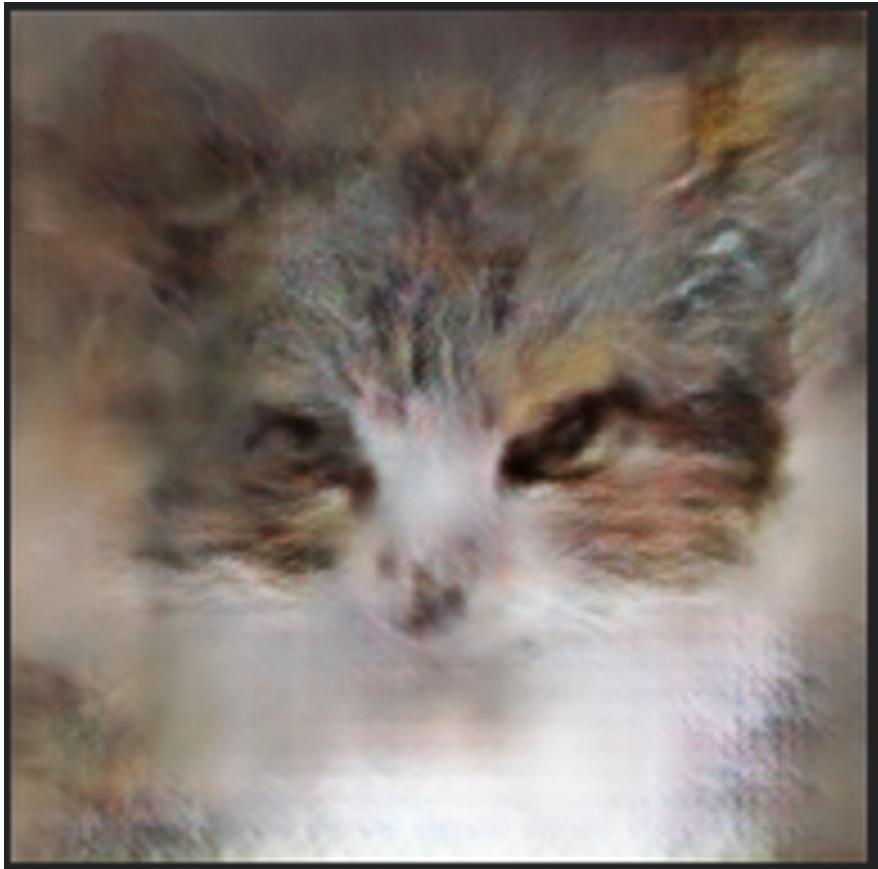
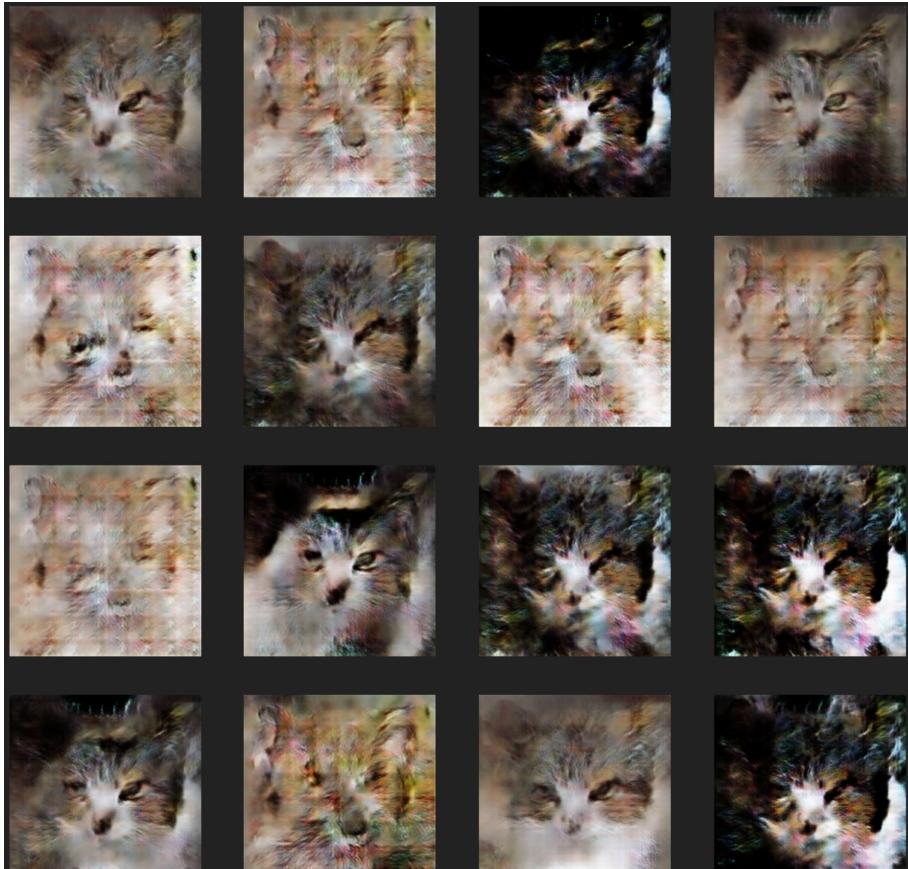
Cats Results - Output Images (500 images) 1k epoch



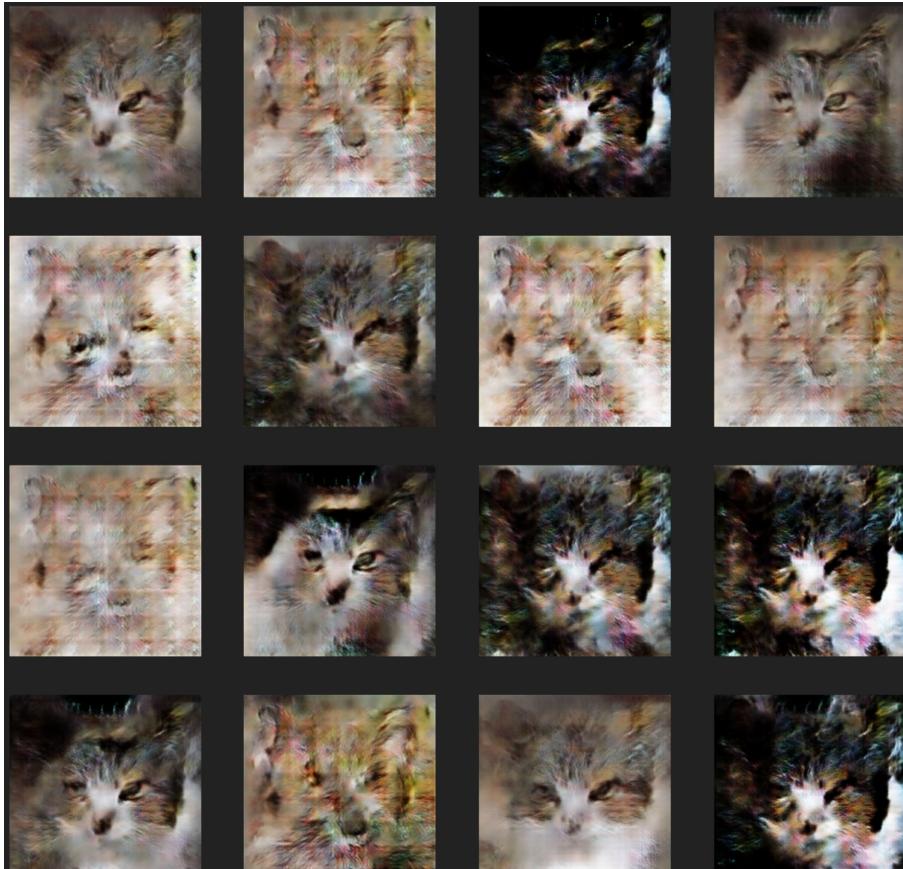
Cats Results - Graphs (500 images) 1k epoch



Cats Results - Output Images (4096 images) 1k epoch

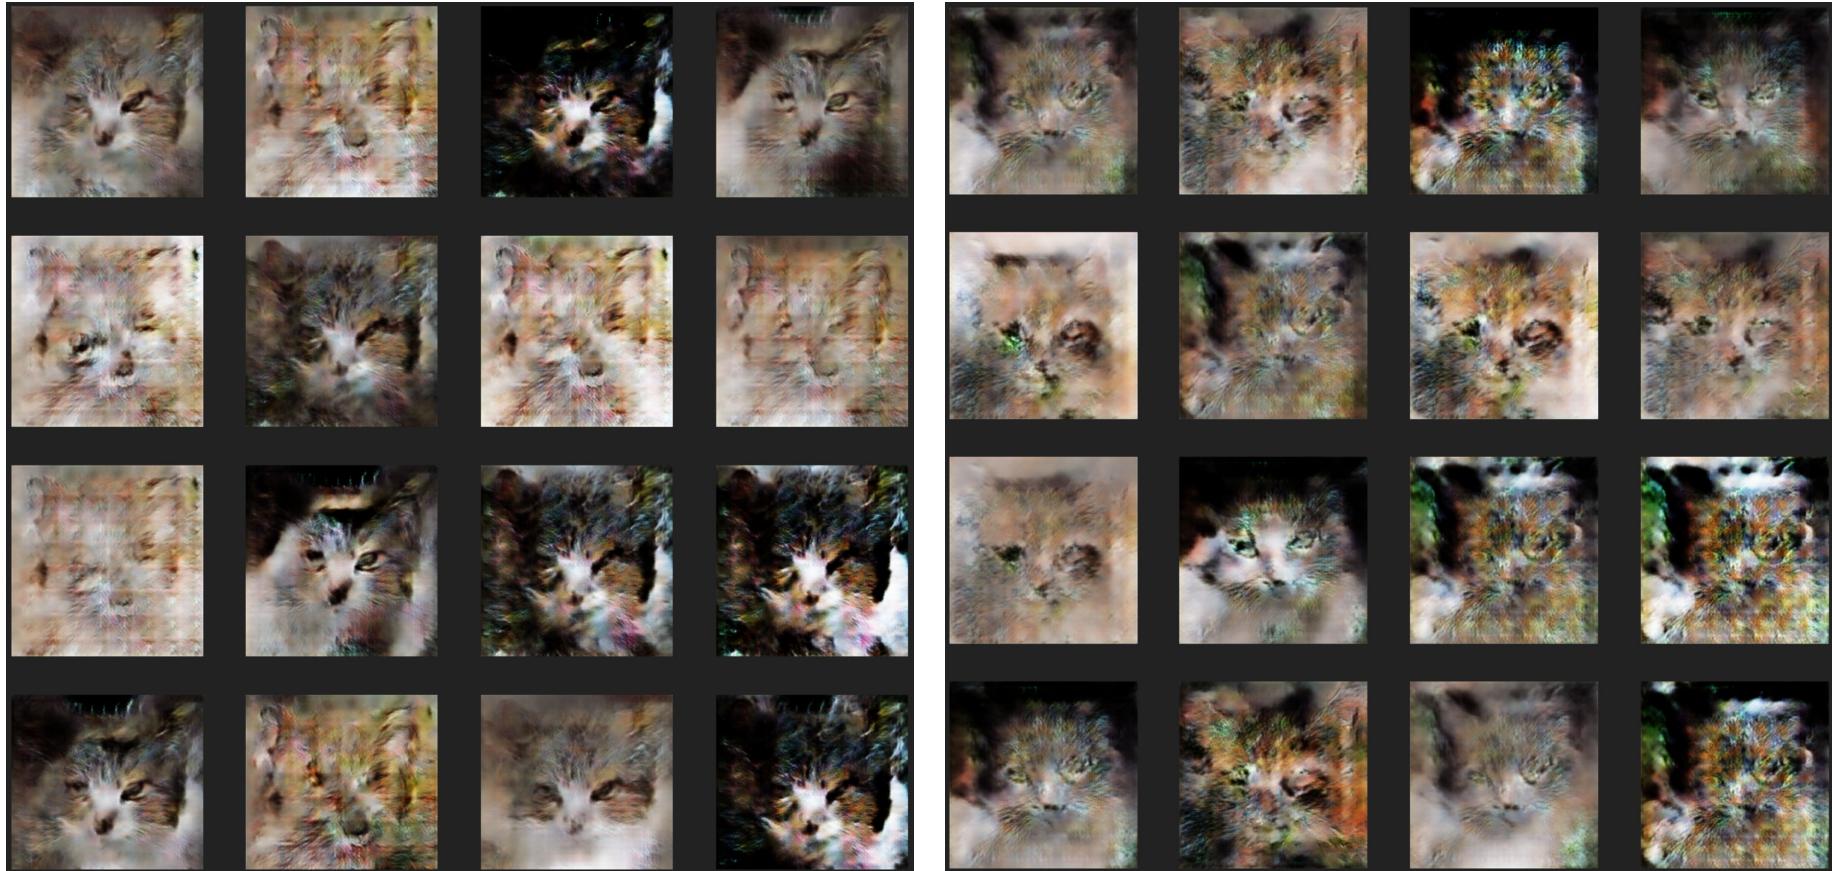


Cats Results - Output Images (4096 images) 1k epoch



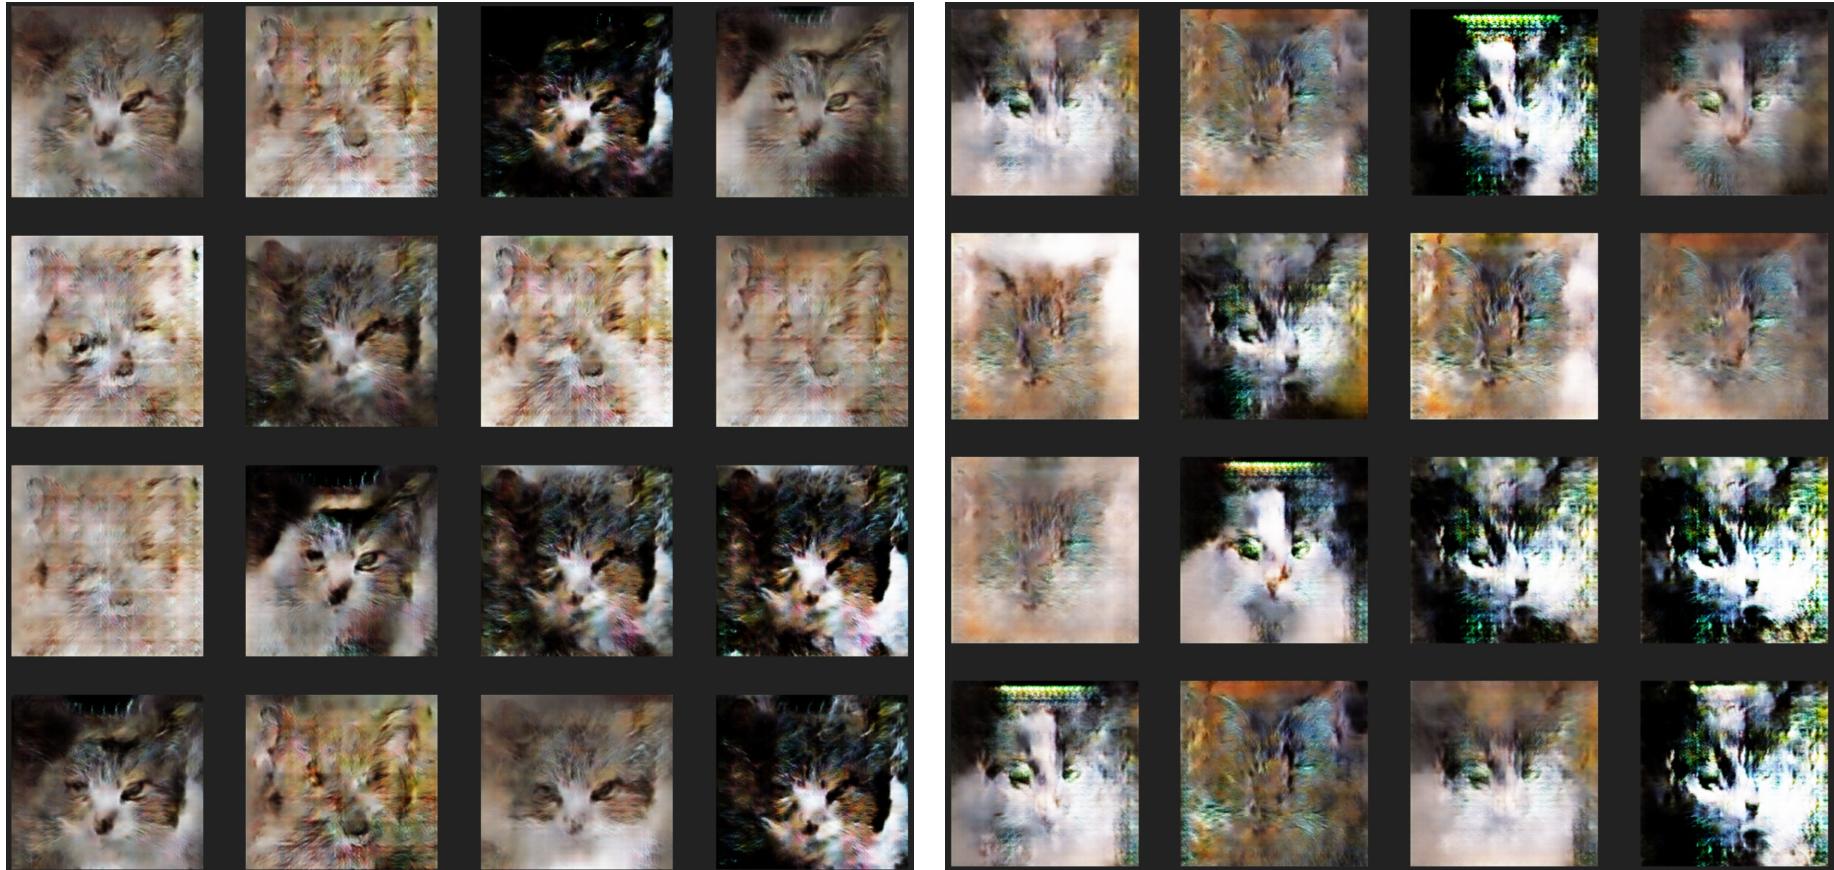
Cats Results - Output Images (4096 images) 1k epoch

900 epoch



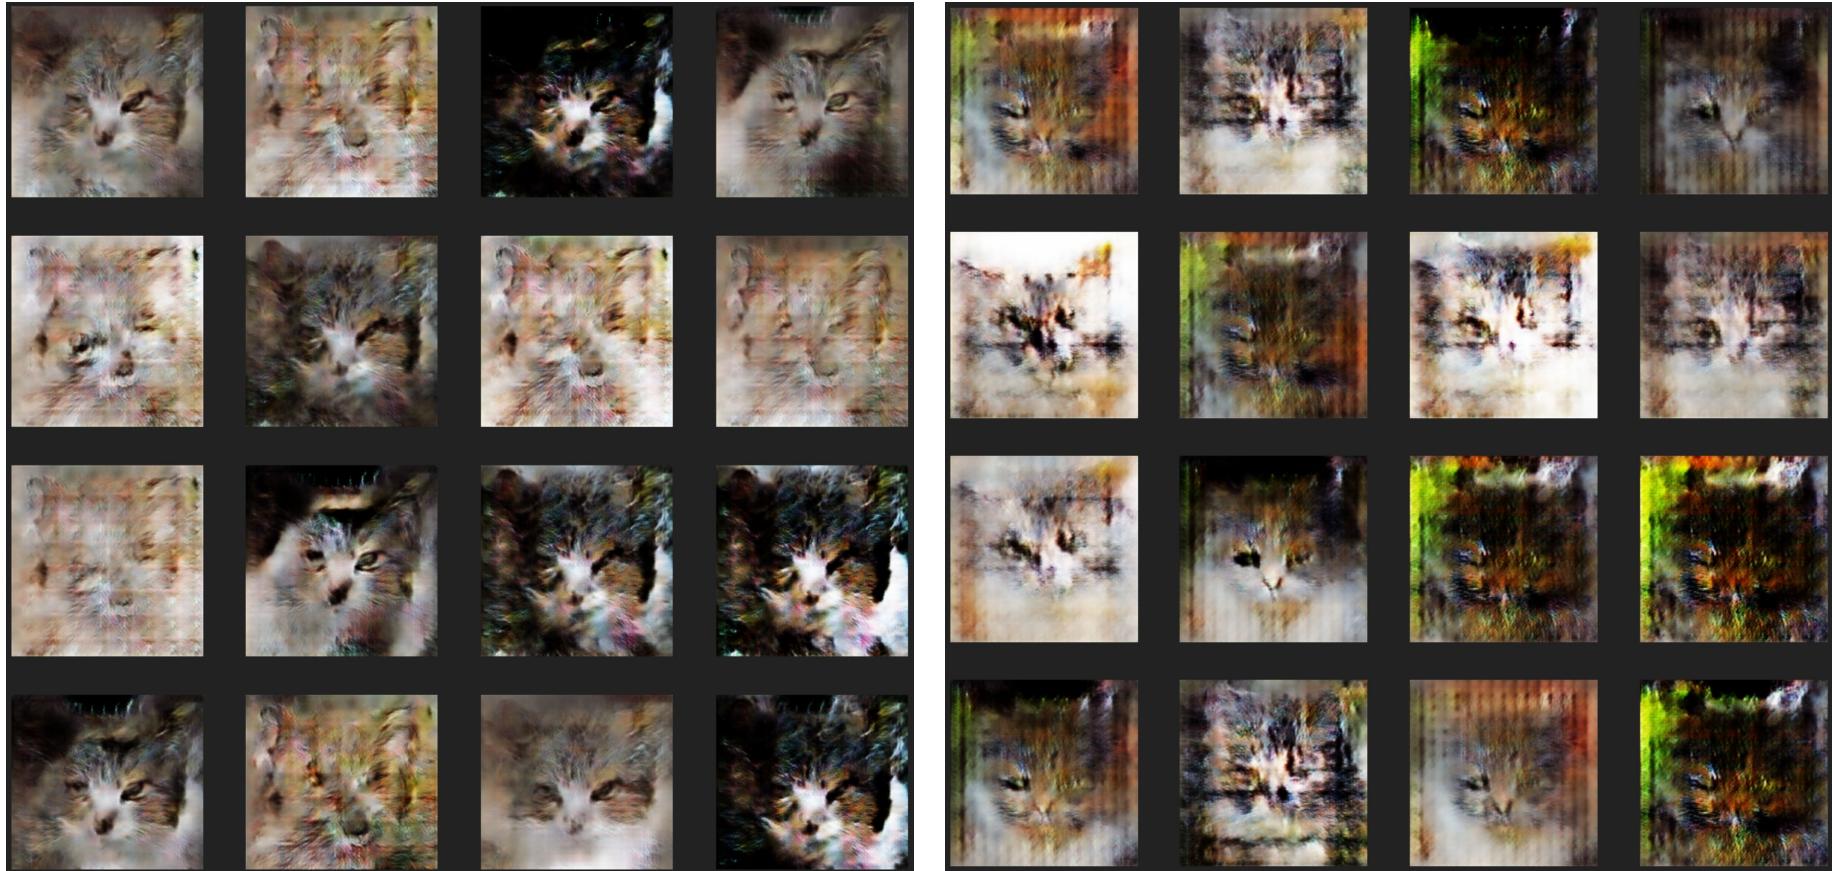
Cats Results - Output Images (4096 images) 1k epoch

800 epoch

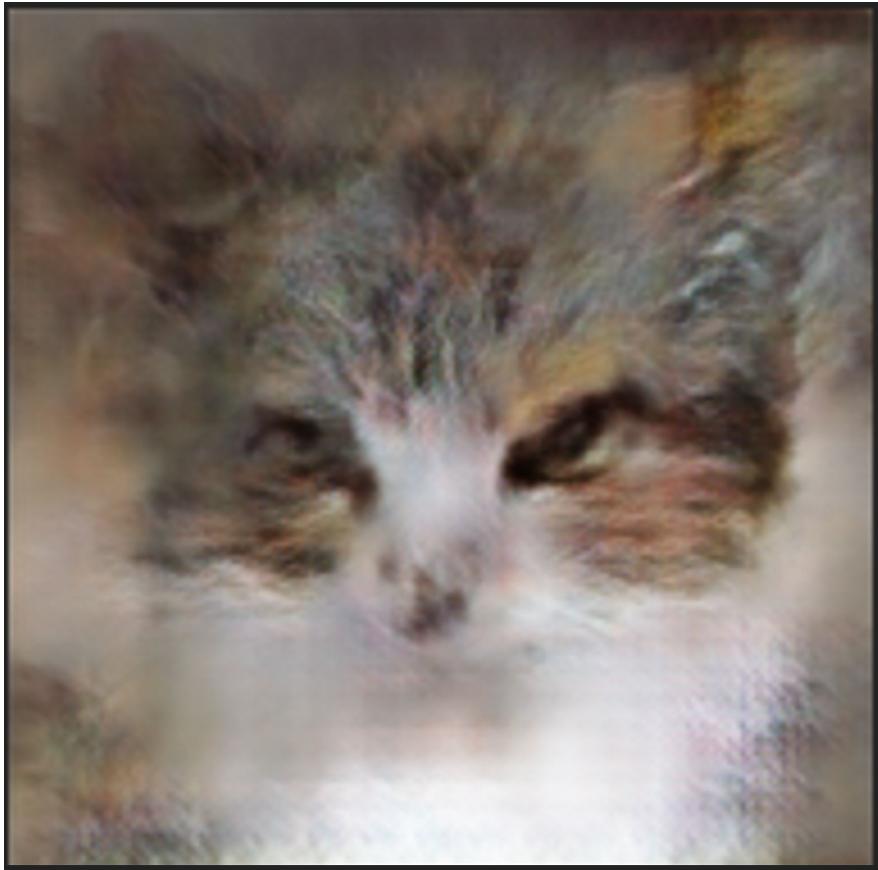


Cats Results - Output Images (4096 images) 1k epoch

700 epoch



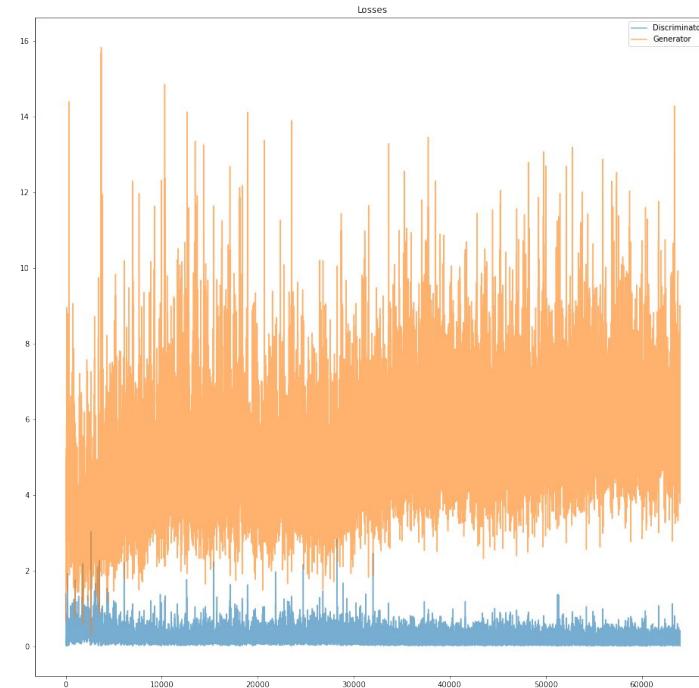
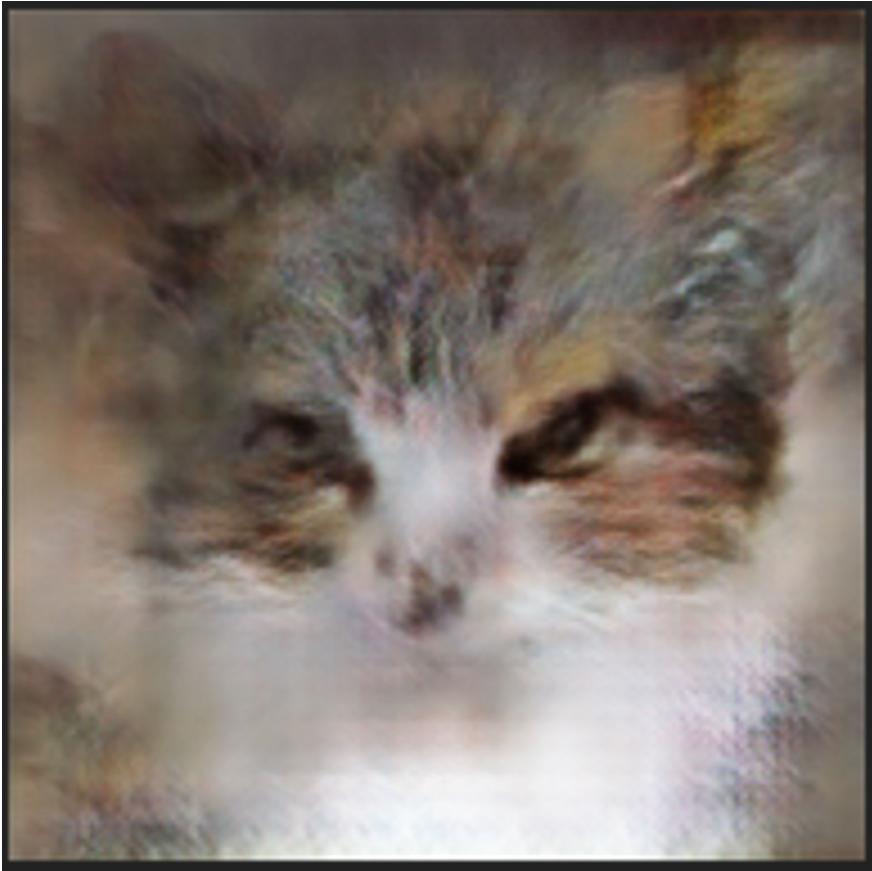
Cats Results - Analyzing Images (4096 images) 1k epoch



Cats Results - Analyzing Images (4096 images) 1k epoch



Cats Results - Graphs (4096 images) 1k epoch



Conclusion of this adventure

- Stick to simple images
- Simpler images = better results sooner
- Avoid images rich in detail, such as 3d images
- Around 1000 images or more
- More images = more epoch
- Diverse pictures work great, such as dogs, cats, cars, horses

Complex images result in poor performance, simpler images results in better performance sooner. Will need to apply different baseline models to see if this still holds true.

Data Preparation

- Finding the data
- Find the “right” data and enough of it
- For example sea shells (shapes and texture)
- Sometimes the data fails to import because one file is bad

Statistics

If your interested:

- Ryzen 7 2700X with 16GB, RTX 3080 Aorus Master, 1TB NVMe
 - FYI, the time per epoch goes up the more we train, thus the range
- Lego, 800 images, 2.2s-2.8s per epoch, ran 5k epoch in 2.5-3 hours
- Cutlery, 870 images, 2.2s-2.8s per epoch, ran 5k epoch in 2.5-3 hours
- Marble, 120 images, 1.2s-1.45s per epoch, ran 5k epoch in 1.5 hours
- Cat, 500 images, 2.3s-3.0s per epoch, ran 1k epoch in 1.5 hours
- Cat, 4096 images, 7.8s-8.6s per epoch, ran 1k in 5 hours
- Have lots of RAM, more than 16GB, with 16GB max ingestion was 4k images.
- Enable GPU with Tensorflow (www.youtube.com/watch?v=lubEtS2JAIY)

References

Datasets

- Cat Dataset: <https://www.kaggle.com/andrewmvd/animal-faces>
- Marble Textures: <https://www.robots.ox.ac.uk/~vgg/data/dtd/>
- Lego Bricks: <https://www.kaggle.com/joosthazelzet/lego-brick-images>
- Cutlery: <https://homepages.inf.ed.ac.uk/rbf/UTENSILS/>

GAN Tutorials and Help

- Tensorflow Tutorial: <https://www.tensorflow.org/tutorials/generative/dcgan>
- Machine Learning Mastery:
<https://machinelearningmastery.com/how-to-develop-a-generative-adversarial-network-for-a-cifar-10-small-object-photographs-from-scratch/>
- Truly Artificial:
<https://simran-tinani.medium.com/truly-artificial-generating-and-experimenting-with-artificial-flower-images-using-deep-d11808193e3b>
- Towards Data Science:
<https://towardsdatascience.com/image-generation-in-10-minutes-with-generative-adversarial-networks-c2afc56bfa3b>